



Workshop on

Domain-Specific Languages for Performance-Portable Weather and Climate Models

Content:

Advanced Concepts II

(Decorations, regional computation, and the future of GT4Py)

Presenter:

Johann Dahm

Learning goals for this session

- Understand the power of the syntax `@gtscript.stencil`
- How to restrict/specialize computation on portions of the IJ axes
- Future design of the toolchain: Storages, horizontal reductions, etc.

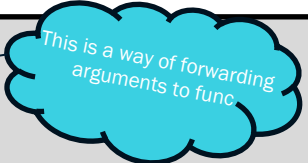
What does @gtscript.stencil actually do?

Python Decorators

A decorator is a function that accepts and returns a function.

```
import time
def timeit(func):
    def wrapper(*args, **kwargs):
        start = time.perf_counter()
        func(*args, **kwargs)
        print("Time: " + str(time.perf_counter() - start))
    return wrapper

@timeit
def say_whee():
    time.sleep(0.5)
    print("Whee!")
```



```
>>> say_whee()
```

What is output when say_whee is executed?

Python Decorators

A decorator is a function that accepts and returns a function.

```
import time
def timeit(func):
    def wrapper(*args, **kwargs):
        start = time.perf_counter()
        func(*args, **kwargs)
        print("Time: " + str(time.perf_counter() - start))
    return wrapper

@timeit
def say_whee():
    time.sleep(0.5)
    print("Whee!")
```

```
>>> say_whee()
Whee!
Time: 0.5012413430013112
```

This is the normal syntax for
decorators

Python Decorators


A decorator is a function that accepts and returns a function.

```
import time
def timeit(func):
    def wrapper(*args, **kwargs):
        start = time.perf_counter()
        func(*args, **kwargs)
        print("Time: " + str(time.perf_counter() - start))
    return wrapper

def say_whee():
    time.sleep(0.5)
    print("Whee!")
say_whee = timeit(say_whee)
```

Since decorators are just functions, they can be called in this way as well with the same effect

How does this relate to GT4Py?



`@gtscript.stencil` returns an entirely new function of generated code instead of a wrapper as `@timeit` does. It:

1. Analyzes the source code from the decorated function
2. Lowers through a series of intermediate representations
3. Generates optimized source code in some language (and if need be, compiles it)
4. Wraps the optimized version in a python function and returns it

Regional Computation

Porting Conditionals?

```
389
390 do j=js-1,jep1
391 do i=is-1,iep1
392 if ( ua(i,j) > 0. ) then
393     if ( i==1 ) then
394         ke(1,j) = uc(1, j)*sin_sg(1,j,1)+v(1,j)*cos_sg(1,j,1)
395     elseif ( i==npx ) then
396         ke(i,j) = uc(npx,j)*sin_sg(npx,j,1)+v(npx,j)*cos_sg(npx,j,1)
397     else
398         ke(i,j) = uc(i,j)
399     endif
400 else
401     if ( i==0 ) then
402         ke(0,j) = uc(1, j)*sin_sg(0,j,3)+v(1,j)*cos_sg(0,j,3)
403     elseif ( i==(npx-1) ) then
404         ke(i,j) = uc(npx,j)*sin_sg(npx-1,j,3)+v(npx,j)*cos_sg(npx-1,j,3)
405     else
406         ke(i,j) = uc(i+1,j)
407     endif
408 endif
409 enddo
410 enddo
```

Conditionals on the iteration position cannot be expressed based on gtscript introduced so far...

Review of the GridTools Parallel Model

Stencils iterate sequentially over **computations** in the order they appear in the code

A computation is composed of **vertical intervals** that are executed sequentially in the order defined by the **iteration policy** of the computation

A vertical interval is executed as a sequential for-loop over the K-range, following the order defined in the iteration policy

Vertical intervals are composed of **statements**, with **horizontal extents** determined automatically

- Parallel model (left) iterates a single statement over the entire horizontal plane
- There are cases where special computation needs to happen on points of the domain
 - BCs
 - Corner cases on multiblock meshes
 - ...

Review of the GridTools Parallel Model

Stencils iterate sequentially over **computations** in the order they appear in the code

A computation is composed of **vertical intervals** that are executed sequentially in the order defined by the **iteration policy** of the computation

A vertical interval is executed as a sequential for-loop over the K-range, following the order defined in the iteration policy

Vertical intervals are composed of **statements**, with **horizontal extents** determined automatically

- Parallel model (left) iterates a single statement over the entire horizontal plane
- There are cases where special computation needs to happen on points of the domain
 - BCs
 - Corner cases on multiblock meshes
 - ...

New Feature: Statements can now be restricted to a portion of the **extended** IJ plane


Frontend Feature

Simple stencil to set an inout field to the following values

```
@gtscript.stencil(...)  
def set_field(field: Field[float]):  
    with computation(PARALLEL), interval(...):  
        field = 0.  
        with parallel(region[I[0]:I[2], :],  
                      region[:, I[-2]:]):  
            field = 1.  
        with parallel(region[:, I[-2]:]):  
            field = 2.
```

The horizontal axes are
called "parallel" axes

Wrap statements in with `parallel(...)`
to restrict to portion(s) of the IJ axes



Frontend Feature

Simple stencil to set an inout field to the following values

```
@gtscript.stencil(...)  
def set_field(field: Field[float]):  
    with computation(PARALLEL), interval(...):  
        field = 0.  
        with parallel(region[I[0]:I[2], :],  
                      region[:, I[-2]:]):  
            field = 1.  
        with parallel(region[:, I[-2]:], I[-2]:):  
            field = 2.
```

The indices used in these slices are relative to the start or end of the *stencil compute domain*

Each argument to `parallel(...)` is a *region*, which is like a 2D array using numpy slice notation

Frontend Feature

Simple stencil to set an inout field to the following values

```
@gtscript.stencil(...)  
def set_field(field: Field[float]):  
    with computation(PARALLEL), interval(...):  
        field = 0.  
        with parallel(region[I[0]:I[2], :],  
                      region[:, I[-2]:]):  
            field = 1.  
        with parallel(region[:, I[-2]:]):  
            field = 2.
```

Statements are still executed in order, so the order matters to get the the corner value correct

Axis Offsets

New **Axis Offset** syntax subscript the axes I, J, K inside the compute domain with wrap-around:

- `region[I[0], :]`: first I-axis point in compute domain (west edge)
- `region[I[-1], :]`: last I-axis point in the compute domain (east edge)

To access points in the **extended compute domain** for a statement, offset these:

- `region[I[0]-1, :]`: point in extended compute domain (if needed)
- `region[I[-1]+1, :]`: point in extended compute domain (if needed)

The region will only exist in generated code if

1. It sets a field value in the compute domain
2. It sets a temporary value that is consumed in the stencil domain (via offset)

Back to the Example...

Scenario: we want to fill a halo value on the east side, then consume it by using that field with an offset...

```
@gtscript.stencil(...)  
def set_field(in_field: Field[float], out_field: Field[float]):  
    with computation(PARALLEL), interval(...):  
        with parallel(region[I[-1] + 1, :]):  
            in_field = 0.5 * (in_field[-1, 0, 0] + in_field[-2, 0, 0])  
            out_field = in_field[1, 0, 0]
```

The value is consumed here, so gt4py generates code for the region

Domain Decomposition & Distributed Memory

Regional computations can also be used to enable domain decomposition in GT4Py

Axis Offsets like `I[0]` can be named externals, such as `istart`

Ranks may have a different value for these externals – leading to different code!

If `istart` is `None`, then regions using this external are automatically removed from the generated code

- `istart = I[0] if rank % 2 else None`
- `jstart = J[0] if rank < 2 else None`
- ...

Domain distributed over 4 ranks

2	3
0	1

Domain Decomposition & Distributed Memory

Regional computations can also be used to enable domain decomposition in GT4Py

Axis Offsets like `I[0]` can be named externals, such as `istart`

Ranks may have a different value for these externals – leading to different code!

If `istart` is `None`, then regions using this external are automatically removed from the generated code

- `istart = I[0] if rank % 2 else None`
- `jstart = J[0] if rank < 2 else None`
- ...

Domain distributed over 4 ranks

2	3
0	1

```
def stencil(...):  
    from __externals__ import istart  
    with computation(PARALLEL), interval(...):  
        with parallel(region[istart, :]):  
            field = 0.
```


Applied to the Example...

```
iend = I[-1] if proc % 2 == 1 else None
@gtscript.stencil(...)
def set_field(in_field: Field[float], out_field: Field[float]):
    from __externals__ import iend
    with computation(PARALLEL), interval(...):
        with parallel(region[iend + 1, :]):
            in_field = 0.5 * (in_field[-1, 0, 0] + in_field[-2, 0, 0])
        out_field = in_field[1, 0, 0]
```

Can offset the Axis Offset in the region slices

FV3 Application: Corner Computation

- Finite difference operators must be specialized at the corners of the cubed sphere tiles, since the values do not really exist
- The hands-on tutorial will walk you through an example that applies the bi-Laplacian operator in a cubed sphere context, where special computation happens at the corners



**Which concepts / motifs that
are missing in GT4Py?**

GT4Py: Stable, yet Active Project

- Formal GDP-based process to introduce new features
- Everyday focus on improving testing and error messages

October 15, 2020 – November 15, 2020

Period: 1 month ▾

Overview

23 Active Pull Requests

12 Active Issues

14

Merged Pull Requests

9

Open Pull Requests


2

Closed Issues

10

New Issues

Excluding merges, **8 authors** have pushed **11 commits** to master and **16 commits** to all branches. On master, **0 files** have changed and there have been **0 additions** and **0 deletions**.



GT4Py Upcoming Features

Recent workshop identified the following items:

- Storages are being replaced by a framework independent of numpy, allowing for cupy, xarray, etc. to be tightly integrated
- Frontend feature to explicitly obtain the I, J, K positions
- Enable compact access syntax compatible with lower-dimensional storages
 - `field[0, 0, 1] = field[K+1]`
 - `field[1, 1, 0] = field[I+1,J+1]`
- Further reduce the overhead of calling a python stencil
- Simple optimization passes to compute redundant code to remove memory reads

[October 2020 GTScript Workshop](#)

GT4Py Future Extensions and Features

We know there are other motifs that may not map directly to stencils that also need to be fast

- Horizontal reductions on fields
- ...

We need your help to identify these. GT4Py can only be great if we all work together to bring ideas from the science domain to the DSL!

Questions?

Hands-on Session

Now it's your turn!

`Session-2B.ipynb` to work on

- Writing an efficient transient diffusion code (§1)
- Exploring corner computation (§2)
- Experimenting with decorators (§3)

See you on Slack! **Next huddle at 4:30pm EST**